

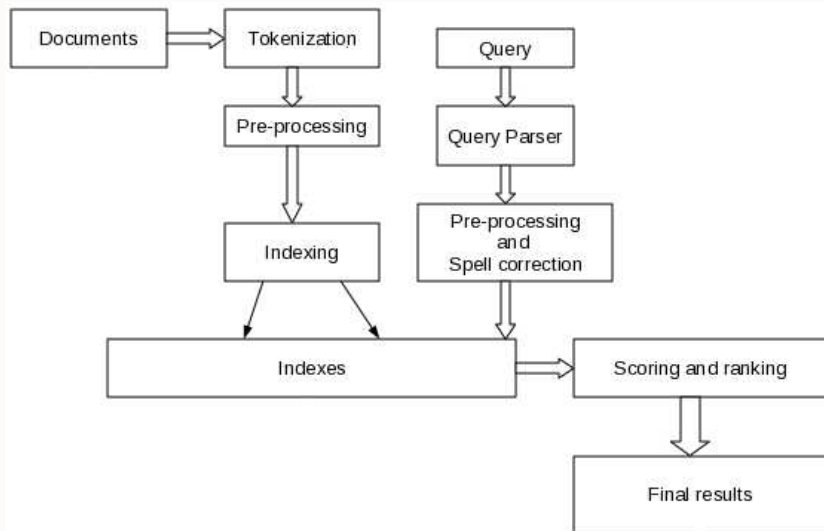
Indexing and Vector Space Model

Jiaul Paik
Indian Statistical Institute , Kolkata

Plan of the Lecture

1. Tokenization
2. Pre-processing
3. Data Structures for Storing Dictionary
4. Index Constructon
5. Term Weighting
6. Vector Space Model

IR system architecture : Overview



Basic indexing steps

- ▶ Collect the documents.
- ▶ Tokenize the input text.
- ▶ Do pre-processing (stopwords, stemming etc.) of tokens.
- ▶ Construct the inverted index.

Tokenization

- ▶ General strategy : chop on white spaces and throw away punctuation characters.
- ▶ Input: “Friends, Romans and Countrymen”.
- ▶ Output: Tokens
 - ▶ Friends
 - ▶ Romans
 - ▶ and
 - ▶ Contrymen
- ▶ A token is an instance of a sequence of characters.
- ▶ Each such token is now a candidate for an index entry, after further processing.

Tokenization : Issues

- ▶ **O'Neill** → neill, oneill, o'neill ?
- ▶ **Hewlett-Packard** → Hewlett and Packard as two tokens?
- ▶ lower-case, lowercase or lower case?
- ▶ San Francisco : One token or two?

Tokenization : language issues

- ▶ Compound words.
 - ▶ German IR benefits significantly (15%) from compound splitting.
 - ▶ Compound nouns without space : *Computerlinguistik* (computational linguistic)
 - ▶ Solution : Subdivide a word into multiple words that appear in vocabulary.


Tokenization : language issues

- ▶ Chinese and japanese has no spaces between words.
- ▶ Arabic (or Hebrew) is basically written right to left, but with certain items like numbers written left to right.

So far ... and now



So far . . . and now



Pre-processing

- ▶ Case folding
- ▶ Stopwords
- ▶ Spelling correction and normalization
- ▶ Stemming
- ▶ Lemmatization

Case folding

- ▶ Reduce all letters to lower case.
- ▶ exceptions
 - ▶ General Motors.
 - ▶ SAIL vs. sail

Stopwords

- ▶ Remove the words/terms from the index which have no or negligible information
- ▶ Pros : Disk space save. (25% to 30%)
- ▶ Cons : Not good if the system has to handle phrase query, song names, etc.
 - ▶ As we may think.
 - ▶ To be or not to be.
- ▶ Common approach
 - ▶ Sort the terms based on descending collection frequency.
 - ▶ Take top few terms.
 - ▶ Hand filter the list.

Spelling normalization and correction

- ▶ **Normalization**

- ▶ color vs. colour.

- ▶ **Correction**

- ▶ Isolated term : attempts to correct single query term at a time.
 - ▶ Edit distance
 - ▶ k-gram overlap
 - ▶ Context sensitive : *flew **form** Heathrow*
 - ▶ Using query log : frequent query

Normalizing morphological variants

- ▶ Inflectional morphology
- ▶ Derivational morphology
- ▶ Two methods : lemmatization and stemming
- ▶ Why we need to care?
 - ▶ Significant performance improvement for more complex languages.
 - ▶ Retrieves more relevant documents.
 - ▶ Index size reduction.

Lemmatization

- ▶ Reduce inflectional/variant forms to base form

Lemmatization

- ▶ Reduce inflectional/variant forms to base form
- ▶ Example: *am, are, is* → *be*

Lemmatization

- ▶ Reduce inflectional/variant forms to base form
- ▶ Example: *am, are, is* \rightarrow *be*
- ▶ Example: *car, cars, car's, cars'* \rightarrow *car*

Lemmatization

- ▶ Reduce inflectional/variant forms to base form
- ▶ Example: *am, are, is* → *be*
- ▶ Example: *car, cars, car's, cars'* → *car*
- ▶ Example: *the boy's cars are different colors* → *the boy car be different color*

Lemmatization

- ▶ Reduce inflectional/variant forms to base form
- ▶ Example: *am, are, is* → *be*
- ▶ Example: *car, cars, car's, cars'* → *car*
- ▶ Example: *the boy's cars are different colors* → *the boy car be different color*
- ▶ Lemmatization implies doing “proper” reduction to dictionary headword form (the **lemma**).

Lemmatization

- ▶ Reduce inflectional/variant forms to base form
- ▶ Example: *am, are, is* → *be*
- ▶ Example: *car, cars, car's, cars'* → *car*
- ▶ Example: *the boy's cars are different colors* → *the boy car be different color*
- ▶ Lemmatization implies doing “proper” reduction to dictionary headword form (the **lemma**).
- ▶ Inflectional morphology (*cutting* → *cut*) vs. derivational morphology (*destruction* → *destroy*)

Stemming

- ▶ Definition of stemming: Crude heuristic process that chops off the ends of words in the hope of achieving what “principled” lemmatization attempts to do with a lot of linguistic knowledge.

Stemming

- ▶ Definition of stemming: Crude heuristic process that chops off the ends of words in the hope of achieving what “principled” lemmatization attempts to do with a lot of linguistic knowledge.
- ▶ Language dependent

Stemming

- ▶ Definition of stemming: Crude heuristic process that chops off the ends of words in the hope of achieving what “principled” lemmatization attempts to do with a lot of linguistic knowledge.
- ▶ Language dependent
- ▶ Often inflectional and derivational

Stemming

- ▶ Definition of stemming: Crude heuristic process that chops off the ends of words in the hope of achieving what “principled” lemmatization attempts to do with a lot of linguistic knowledge.
- ▶ Language dependent
- ▶ Often inflectional and derivational
- ▶ Example for derivational: *automate*, *automatic*, *automation* all reduce to *automat*

Stemming

- ▶ Definition of stemming: Crude heuristic process that chops off the ends of words in the hope of achieving what “principled” lemmatization attempts to do with a lot of linguistic knowledge.
- ▶ Language dependent
- ▶ Often inflectional and derivational
- ▶ Example for derivational: *automate*, *automatic*, *automation* all reduce to *automat*
- ▶ Common strategy : Identify (manually or statistically) a set of suffixes and remove from the ends of words

Porter algorithm

- ▶ Most common algorithm for stemming English

Porter algorithm

- ▶ Most common algorithm for stemming English
- ▶ Results suggest that it is at least as good as other stemming options

Porter algorithm

- ▶ Most common algorithm for stemming English
- ▶ Results suggest that it is at least as good as other stemming options
- ▶ 5 phases of reductions

Porter algorithm

- ▶ Most common algorithm for stemming English
- ▶ Results suggest that it is at least as good as other stemming options
- ▶ 5 phases of reductions
- ▶ Phases are applied sequentially

Porter algorithm

- ▶ Most common algorithm for stemming English
- ▶ Results suggest that it is at least as good as other stemming options
- ▶ 5 phases of reductions
- ▶ Phases are applied sequentially
- ▶ Each phase consists of a set of commands.

Porter algorithm

- ▶ Most common algorithm for stemming English
- ▶ Results suggest that it is at least as good as other stemming options
- ▶ 5 phases of reductions
- ▶ Phases are applied sequentially
- ▶ Each phase consists of a set of commands.
 - ▶ Sample command: Delete final *ement* if what remains is longer than 1 character

Porter algorithm

- ▶ Most common algorithm for stemming English
- ▶ Results suggest that it is at least as good as other stemming options
- ▶ 5 phases of reductions
- ▶ Phases are applied sequentially
- ▶ Each phase consists of a set of commands.
 - ▶ Sample command: Delete final *ement* if what remains is longer than 1 character
 - ▶ replacement → replac

Porter algorithm

- ▶ Most common algorithm for stemming English
- ▶ Results suggest that it is at least as good as other stemming options
- ▶ 5 phases of reductions
- ▶ Phases are applied sequentially
- ▶ Each phase consists of a set of commands.
 - ▶ Sample command: Delete final *ement* if what remains is longer than 1 character
 - ▶ replacement → replac
 - ▶ cement → cement

Porter algorithm

- ▶ Most common algorithm for stemming English
- ▶ Results suggest that it is at least as good as other stemming options
- ▶ 5 phases of reductions
- ▶ Phases are applied sequentially
- ▶ Each phase consists of a set of commands.
 - ▶ Sample command: Delete final *ement* if what remains is longer than 1 character
 - ▶ replacement → replac
 - ▶ cement → cement
- ▶ Sample convention: Of the rules in a compound command, select the one that applies to the longest suffix.

Porter stemmer: A few rules

Rule

SSSES → SS

IES → I

SS → SS

S →

Example

caresses → caress

ponies → poni

caress → caress

cats → cat

No stem, stem and lemmatization : A comparison

Input text: Such an analysis can reveal features that are not easily visible from the variations in the individual genes and can lead to a picture of expression that is more biologically transparent and accessible to interpretation

No stem, stem and lemmatization : A comparison

Input text: Such an analysis can reveal features that are not easily visible from the variations in the individual genes and can lead to a picture of expression that is more biologically transparent and accessible to interpretation

Porter stemmer: such an analysi can reveal featur that ar not easili visibl from the variat in the individu gene and can lead to a pictur of express that is more biolog transpar and access to interpret

No stem, stem and lemmatization : A comparison

Input text: Such an analysis can reveal features that are not easily visible from the variations in the individual genes and can lead to a picture of expression that is more biologically transparent and accessible to interpretation


Porter stemmer: such an analysi can reveal featur that ar not easili visibl from the variat in the individu gene and can lead to a pictur of express that is more biolog transpar and access to interpret

Lemmatized text : Such an analysis can reveal feature that are not easily visible from the variation in the individual gene and can lead to a picture of expression that is more biologically transparent and accessible to interpretation


Does stemming improve effectiveness?

- ▶ In general, stemming increases effectiveness for some queries, and decreases effectiveness for others.
- ▶ Some languages (Hungarian, Czech) enjoy 30% to 40% benefits.
- ▶ Queries where stemming is likely to help
 - ▶ tartan sweaters
 - ▶ (equivalence class: {sweater,sweaters})
 - ▶ sightseeing tour san francisco
 - ▶ (equivalence class: {tour,tours})
- ▶ Porter Stemmer maps all of {*operate, operating, operates, operation, operative, operatives, operational*} to *oper*.
- ▶ Queries where stemming hurts:
 - ▶ operational research
 - ▶ operating system
 - ▶ operative dentistry
- ▶ Stemming reduces the vocabulary size significantly.

So far ... and now



So far ... and now



Inverted index

- For each term t , we store a list of all documents that contain t along with the tf .

Input : Documents

id1 : Web mining is useful.

id2 : Usage mining applications.

id3 : Web structure mining studies the hyperlink structure of web.

Output : Inverted Index

Dictionary	Postings
applications	: $\langle id2, 1, [3] \rangle$
hyperlink	: $\langle id3, 1, [6] \rangle$
mining	: $\langle id1, 1, [2] \rangle, \langle id2, 1, [2] \rangle, \langle id3, 1, [3] \rangle$
structure	: $\langle id3, 2, [2, 7] \rangle$
studies	: $\langle id3, 1, [4] \rangle$
usage	: $\langle id2, 1, [1] \rangle$
useful	: $\langle id1, 1, [4] \rangle$
web	: $\langle id1, 1, [1] \rangle, \langle id3, 2, [1, 8] \rangle$

Dictionary data structures for inverted index

- ▶ The dictionary data structure stores the term vocabulary, document frequency, pointers to each postings list.
 - ▶ Term vocabulary : the data
 - ▶ Document frequency : the no. of documents that contain the term

Dictionary as array of fixed-width entries

- ▶ For each term, we need to store a couple of items:
 - ▶ document frequency
 - ▶ pointer to postings list
 - ▶ ...
- ▶ Assume for the time being that we can store this information in a fixed-length entry.
- ▶ Assume that we store these entries in an array.

Dictionary as array of fixed-width entries

term	document frequency	pointer to postings list
a	656,265	→
aabir	65	→
...
zonathan	221	→

Dictionary as array of fixed-width entries

term	document frequency	pointer to postings list
a	656,265	→
aabir	65	→
...
zonathan	221	→

How do we look up a query term q_i in this array at query time? That is: which data structure do we use to locate the entry (row) in the array where q_i is stored?

Data structures for looking up term

- ▶ Two main classes of data structures: hashes and trees
- ▶ Some IR systems use hashes, some use trees.
- ▶ Criteria for when to use hashes vs. trees:
 - ▶ Is there a fixed number of terms or will it keep growing?
 - ▶ How many terms are we likely to have?

Hashes

- ▶ Each vocabulary term is hashed into an integer.
- ▶ Try to avoid collisions
- ▶ At query time, do the following: hash query term, resolve collisions, locate entry in fixed-width array
- ▶ Pros: Lookup in a hash is faster than lookup in a tree.
 - ▶ Lookup time is constant.
- ▶ Cons
 - ▶ no prefix search (all terms starting with *automat*)
 - ▶ need to rehash everything periodically if vocabulary keeps growing

Trees

- ▶ Trees solve the prefix problem (find all terms starting with *automat*).
- ▶ Simplest tree: binary search tree
- ▶ Search is slightly slower than in hashes: $O(\log M)$, where M is the size of the vocabulary.
- ▶ $O(\log M)$ only holds for balanced trees.
- ▶ Rebalancing binary search trees is expensive.
- ▶ B-trees mitigate the rebalancing problem.
- ▶ B-tree definition: every internal node has a number of children in the interval $[a, b]$ where a, b are appropriate positive integers, e.g., $[2, 4]$.
- ▶ Trie

Index Construction

- ▶ How do we construct an index efficiently?
- ▶ What strategies can we use with limited main memory?

Hardware basics

- ▶ Access to data in main memory faster than in disk.
- ▶ Disk seeks: No data is transferred from disk while the disk head is being positioned.
- ▶ Therefore: Transferring one large chunk of data from disk to memory is faster than transferring many small chunks.
- ▶ Disk I/O is blockbased: Reading and writing of entire blocks (as opposed to smaller chunks).

Index Construction : Key Steps

- ▶ Documents are parsed to extract words and these are saved with the Document ID.
- ▶ After all documents have been parsed, the inverted file is sorted by terms.

Scaling Index Construction

- ▶ In-memory index construction does not scale.
- ▶ How can we construct an index for very large collections?

Sort Based Index Construction

- ▶ As we build the index, we parse docs one at a time.
- ▶ The final postings for any term are incomplete until the end.
- ▶ Store the term (or termID), docid pairs in main memory buffer.
- ▶ Sort (term/termID as primary key, docid as secondary key) and write to intermediate files in disk when the buffer is full.
- ▶ Merge all intermediate files into a sorted file.
- ▶ Finally build the inverted index in a linear scan.


Problem with sort based indexing

- ▶ Assumption : we can keep dictionary into main memory.
- ▶ We need the dictionary (which grows dynamically) in order to implement a term to termID mapping.
- ▶ Actually, we could work with term, docID postings instead of termID, docID postings. But then intermediate files are large and we end up with a scalable but slow method.

Single pass in-memory indexing

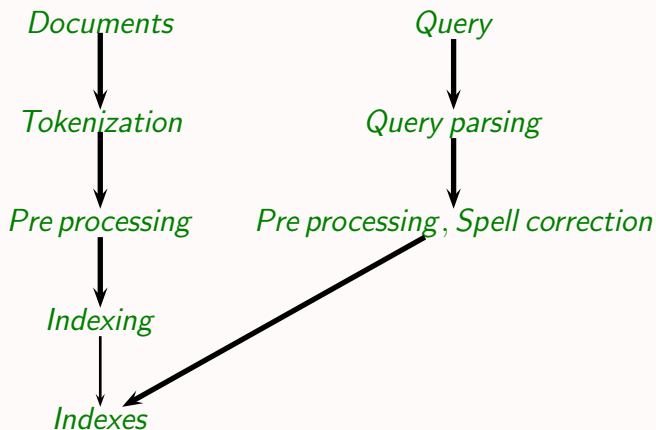
- ▶ Key idea 1: Generate separate dictionaries for each block
no need to maintain term-termID mapping across blocks.
- ▶ Key idea 2: Accumulate postings in postings lists as they occur.
- ▶ With these two ideas we can generate a complete inverted index for each block.
- ▶ These separate indexes can then be merged into one big index.

So far ...

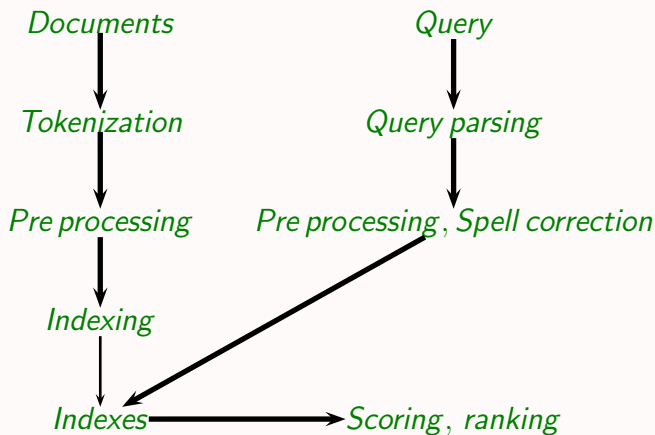


TEA BREAK

So far ... and now



So far ... and now



The retrieval problem

- ▶ Input
 - ▶ A set of documents : $d_1, d_2, \dots d_n$
 - ▶ A query $q = q_1 q_2 \dots q_n$
- ▶ Output : A set of documents relevant to the query q .
- ▶ The first idea
 - ▶ Just look if a document contain the query terms.
 - ▶ Not good because too much / or too few match.
 - ▶ Users hardly read beyond first few.
 - ▶ Solution : Ranked retrieval (return documents based on the matching score)

Ranked retrieval

- ▶ Rather than a set of documents satisfying a query expression, in ranked retrieval models, the system returns an ordering over the (top) documents in the collection with respect to a query
- ▶ **Free text queries:** Rather than a query language of operators and expressions, the user's query is just one or more words in a human language

Scoring as the basis of ranked retrieval

- ▶ We wish to return in order the documents most likely to be useful to the searcher
- ▶ How can we rank the documents in the collection with respect to a query?
- ▶ Assign a score to each document with respect to the query
- ▶ This score measures how well document and query match.

Query-document matching scores

- ▶ We need a way of assigning a score to a query-document pair
- ▶ Let's start with a one-term query
 - ▶ If the query term does not occur in the document: score should be 0
 - ▶ The more frequent the query term in the document, the higher the score (should be)
 - ▶ We will look at a number of alternatives for this.
- ▶ For the rest of the talk : bag-of-words model assumption
 - ▶ Doesn't consider the ordering of words in a document
 - ▶ John is quicker than Mary and Mary is quicker than John have the same representation

Term weighting : term frequency

- ▶ The term frequency $tf_{t,d}$ of a term t in a document d is defined as the no. of times t occurs in document d .
- ▶ We want to use tf as a measure of importance when matching query with a document.
- ▶ Raw term frequency is not a good idea.
 - ▶ A document with 5 occurrences of a term is more relevant than a document with 1 occurrence of the same term.
 - ▶ But not 5 times relevant.
- ▶ Relevance does not increase proportionally with term frequency.

Scaling term frequency

- ▶ The log frequency weight of term t in d is defined as follows

$$w_{t,d} = \begin{cases} 1 + \log_{10} \text{tf}_{t,d} & \text{if } \text{tf}_{t,d} > 0 \\ 0 & \text{otherwise} \end{cases}$$

- ▶ $\text{tf}_{t,d} \rightarrow w_{t,d}$:
 $0 \rightarrow 0, 1 \rightarrow 1, 2 \rightarrow 1.3, 10 \rightarrow 2, 1000 \rightarrow 4$, etc.
- ▶ Score for a document-query pair: sum over terms t in both q and d :
tf-matching-score(\mathbf{q}, \mathbf{d}) = $\sum_{t \in q} (1 + \log \text{tf}_{t,d})$
- ▶ The score is 0 if none of the query terms is present in the document.

Desired weight for rare terms

- ▶ Rare terms are more informative than frequent terms.
- ▶ Consider a term in the query (**capricious person**) that is **rare** in the collection (e.g., **capricious**).
- ▶ A document containing this term is very likely to be relevant.
- ▶ → We want **high weights** for rare terms like **capricious**.

Desired weight for frequent terms

- ▶ Frequent terms are less informative than rare terms.
- ▶ Consider a term in the query that is frequent in the collection (e.g., **good**, **increase**, **person**).
- ▶ A document containing this term is more likely to be relevant than a document that doesn't ...
- ▶ ...but words like **good**, **increase** and **person** are not sure indicators of relevance.
- ▶ → For frequent terms like **good**, **increase**, and **person**, we want positive weights ...
- ▶ ...but lower weights than for rare terms.

Document frequency

- ▶ We want high weights for rare terms like **capricious**.
- ▶ We want low (positive) weights for frequent words like **good**, **increase**, and **person**.
- ▶ We will use document frequency to factor this into computing the matching score.
- ▶ The document frequency is the number of documents in the collection that the term occurs in.

idf weight

- ▶ df_t is the document frequency, the number of documents that t occurs in.
- ▶ df_t is an inverse measure of the **informativeness** of term t .
- ▶ We define the **idf weight** of term t as follows:

$$idf_t = \log_{10} \frac{N}{df_t}$$

(N is the number of documents in the collection.)

- ▶ idf_t is a measure of the **informativeness** of the term.
- ▶ $[\log N/df_t]$ instead of $[N/df_t]$ to “dampen” the effect of idf
- ▶ Note that we use the log transformation for both term frequency and document frequency.

Examples for idf

Compute idf_t using the formula: $\text{idf}_t = \log_{10} \frac{1,000,000}{\text{df}_t}$

term	df_t	idf_t
calpurnia	1	6
animal	100	4
sunday	1000	3
fly	10,000	2
under	100,000	1
the	1,000,000	0

Effect of idf on ranking

- ▶ idf has no effect on ranking for one-term queries.
 - ▶ For example : iPhone
- ▶ idf affects the ranking of documents for queries with at least two terms.
- ▶ For example, in the query *capricious person*, idf weighting increases the relative weight of **capricious** and decreases the relative weight of **person**.

Collection frequency vs. Document frequency

word	collection frequency	document frequency
insurance	10440	3997
try	10422	8760

- ▶ Collection frequency of t : number of tokens of t in the collection
- ▶ Document frequency of t : number of documents t occurs in
- ▶ Which word is a better search term (and should get a higher weight)?
- ▶ This example suggests that df (and idf) is better for weighting than cf (and “ icf ”).

tf-idf weighting

- ▶ The tf-idf weight of a term is the product of its tf weight and its idf weight.
- ▶ $w_{t,d} = (1 + \log \text{tf}_{t,d}) \cdot \log \frac{N}{\text{df}_t}$
- ▶ Best known weighting scheme in information retrieval

Summary: tf-idf

- ▶ Assign a tf-idf weight for each term t in each document d : $w_{t,d} = (1 + \log \text{tf}_{t,d}) \cdot \log \frac{N}{\text{df}_t}$
- ▶ The tf-idf weight ...
 - ▶ Increases with the number of occurrences within a document. (term frequency)
 - ▶ Increases with the rarity of the term in the collection. (inverse document frequency)
 - ▶ Benefit is maximum when a term occurs many times within a small no. of documents
 - ▶ Lowest when the term occurs in virtually all documents.

Scoring function

► $\text{Score}(\mathbf{q}, \mathbf{d}) = \sum_{t \in \mathbf{q}} \mathbf{w}_{t,\mathbf{q}} \times \mathbf{w}_{t,\mathbf{d}}$

For example, $w_{t,q}$ may be no. of time term t occurs in q

Vector Space Model : Documents as vectors

- ▶ Each document is now represented as a real-valued vector of tf-idf weights $\in \mathbb{R}^{|V|}$. $|V|$ is the no. of distinct terms in the collection.
- ▶ So we have a $|V|$ -dimensional real-valued vector space.
- ▶ Terms are **axes** of the space.
- ▶ Documents are **points** or **vectors** in this space.
- ▶ Very high-dimensional: tens of millions of dimensions when you apply this to web search engines
- ▶ Each vector is very sparse - most entries are zero.


Queries as vectors

- ▶ Key idea 1: do the same for queries: represent them as vectors in the high-dimensional space
- ▶ Key idea 2: Rank documents according to their proximity to the query
- ▶ proximity = similarity
- ▶ proximity \approx inverse of distance


How do we formalize vector space similarity?

- ▶ First cut: distance between two points?
- ▶ Euclidean distance?
- ▶ Euclidean distance is a bad idea
 - ▶ Because Euclidean distance is large for vectors of different lengths.

Why distance is a bad idea?



Why distance is a bad idea?



The Euclidean distance of \vec{q} and \vec{d}_2 is large although the distribution of terms in the query q and the distribution of terms in the document d_2 are very similar.

Use angle instead of distance

- ▶ Rank documents according to angle with query
- ▶ An experiment
 - ▶ Take a document d and append it to itself. Call this document d' . d' is twice as long as d .
 - ▶ For example let $d = \textit{Sachin is a cricketer}$ and $d' = \textit{Sachin is a cricketer. Sachin is a cricketer}$
 - ▶ “Semantically” d and d' have the same content.
 - ▶ The angle between the two documents is 0, corresponding to maximal similarity ...
 - ▶ ... even though the Euclidean distance between the two documents can be quite large.

From angles to cosines

- ▶ The following two notions are equivalent.
 - ▶ Rank documents according to the angle between query and document in decreasing order
 - ▶ Rank documents according to cosine (query, document) in increasing order
- ▶ Cosine is a monotonically decreasing function of the angle for the interval $[0^\circ, 180^\circ]$

Ranked retrieval in vector space model : summary

- ▶ Represent the query as a weighted tf-idf vector
- ▶ Represent each document as a weighted tf-idf vector
- ▶ Compute the cosine similarity between the query vector and each document vector
- ▶ Rank documents based on the similarity score
- ▶ Return the top K (e.g., $K = 10$) to the user

Computing the cosine score using inverted index

COSINESCORE(q)

```
1  float Scores[ $N$ ] = 0
2  float Length[ $N$ ]
3  for each query term  $t$ 
4  do calculate  $w_{t,q}$  and fetch postings list for  $t$ 
5      for each pair( $d, tf_{t,d}$ ) in postings list
6      do Scores[ $d$ ] + =  $w_{t,d} \times w_{t,q}$ 
7  Read the array Length
8  for each  $d$ 
9  do Scores[ $d$ ] = Scores[ $d$ ] / Length[ $d$ ]
10 return Top  $K$  components of Scores[]
```

References

1. Introduction to Information Retrieval. 2008. By Manning et. al.
2. Managing Gigabytes: Compressing and Indexing Documents and Images. 1999. By Witten et.al.
3. Term Weighting Approaches to Automatic Text Retrieval. 1998. by Salton and Buckley

THANK YOU!