Data Mining
(CS C 415/CS F 415/IS C415)

BITS Pilani
Pilani | Dubai | Goa | Hyderabad

BITS Pilani
K K Birla Goa Campus

Mrs.Aruna.G
Dept. of CS and IS

Text and Reference Books

Prescribed Text Book (S)

T1. Tan,Pang-Ning and other “Introduction to Data Mining” Pearson Education,2006.

Reference Book (S)

R1. Han J & Kamber M, “**Data Mining: Concepts and Techniques”,** Morgan Kaufmann Publishers, 2001

R2. Hand D, Mannila H, & Smyth P, “***Principles of Data Mining***”, MIT Press, 2001.

R3. Pujari A K, "Data Mining Techniques", University Press (India), 2001.

R4. Kimball R, "The Data Warehouse Toolkit", 2e, John Wiley, 2002.

Evaluation Components

- Test 1 : 20%
- Test 2 : 20%
- Project/Presentation : 30%
- Compre : 30%

Lecture 1: Introduction

Road Map

- Motivation
- What is DM?
- DM Tasks
- Applications
- Issues / Challeneges of DM

Evolution of Database Technology

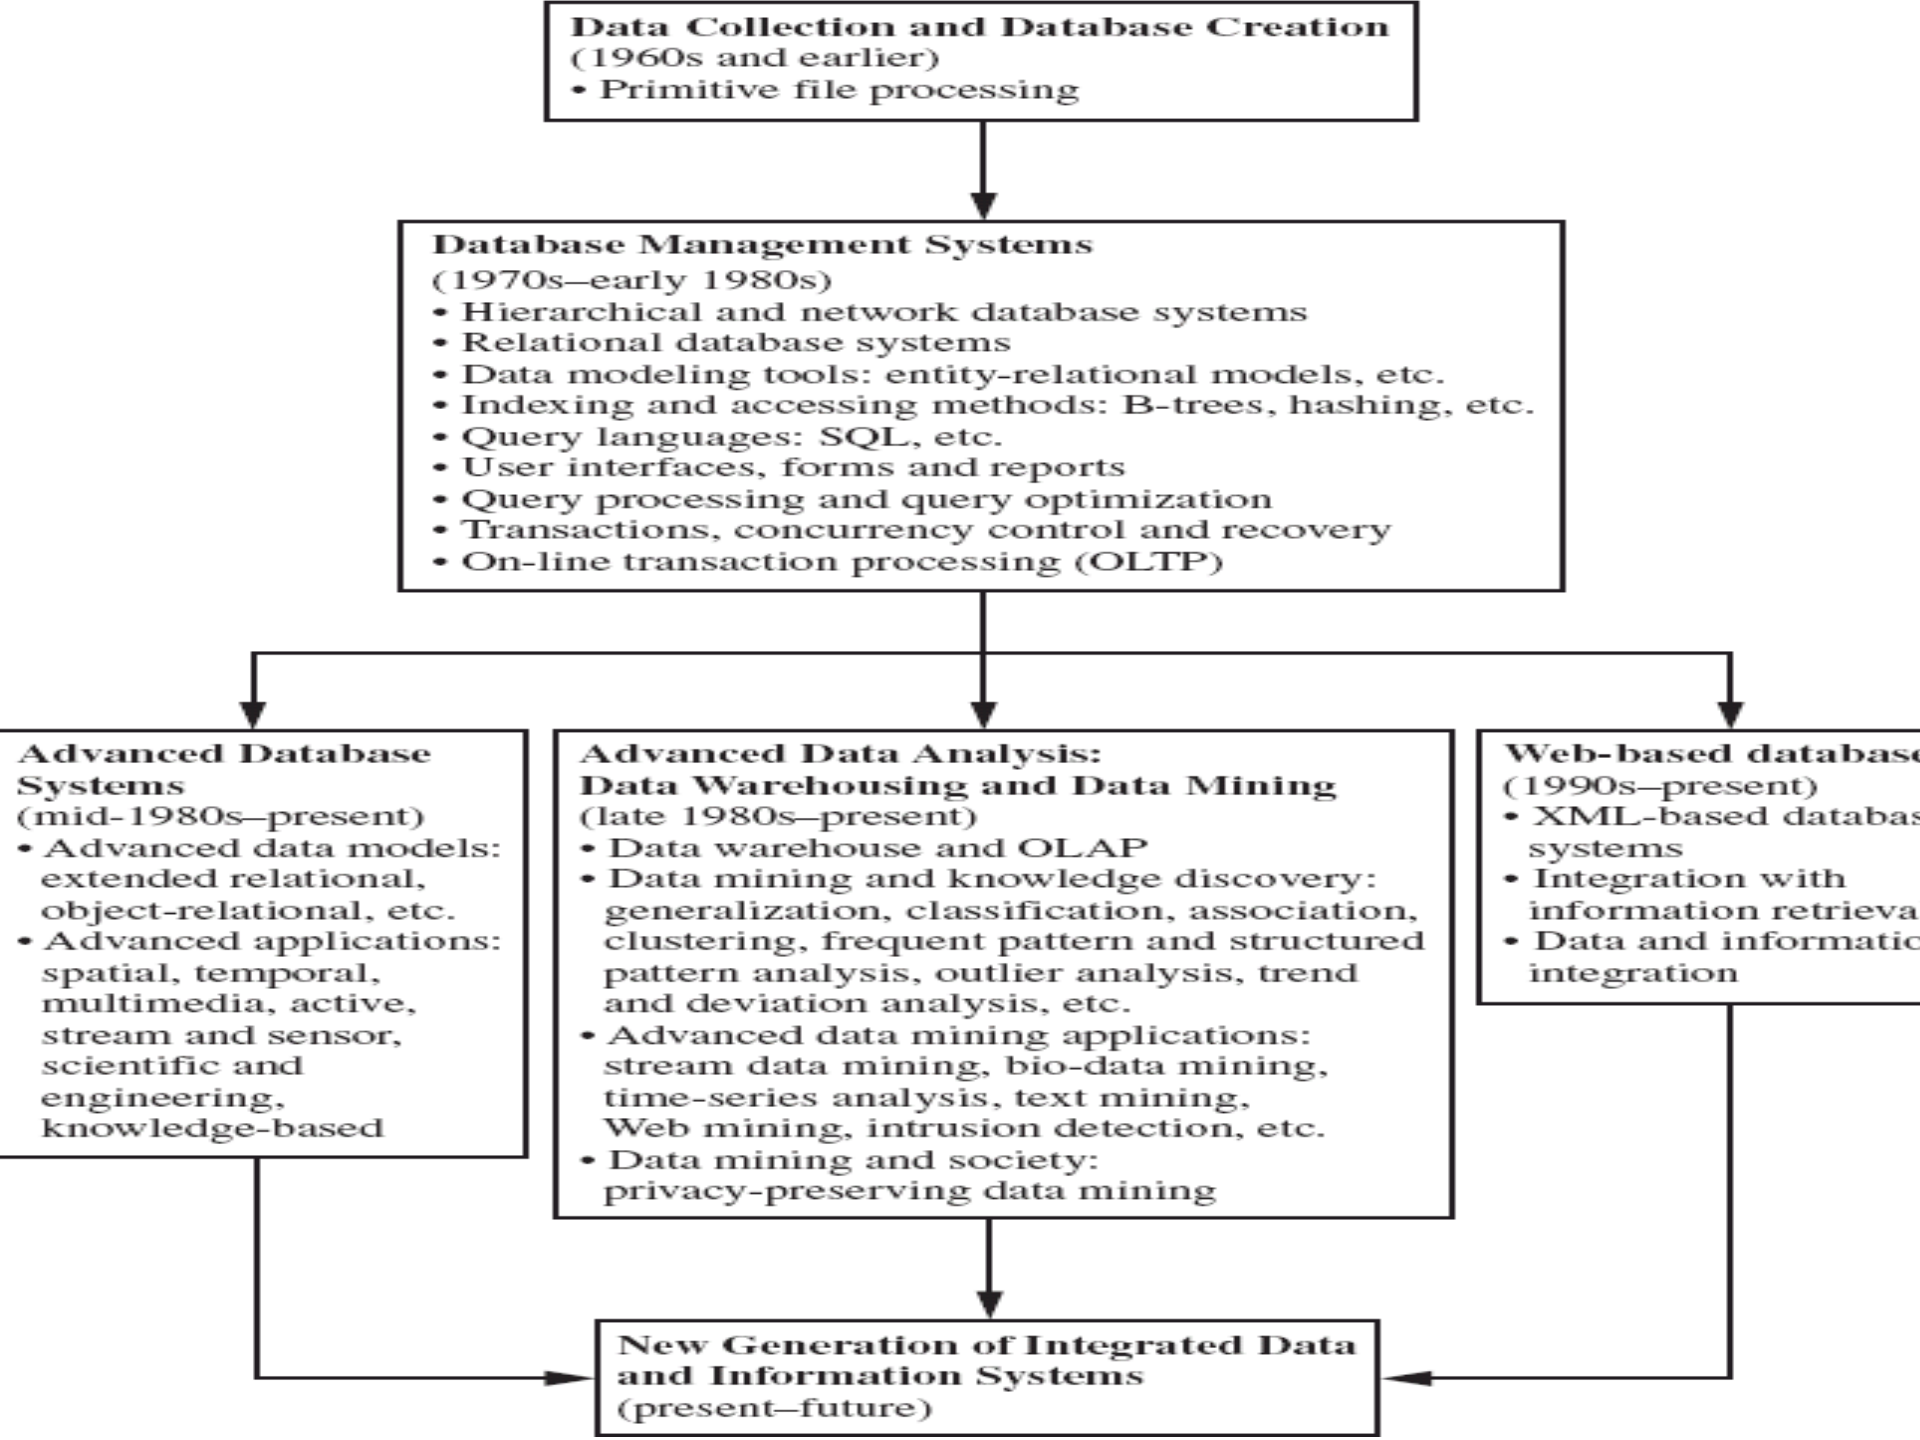
Data Collection and Database Creation
(1960s and earlier)
• Primitive file processing
Database Management Systems
(1970s–early 1980s)
• Hierarchical and network database systems
• Relational database systems
• Data modeling tools: entity-relational models, etc.
• Indexing and accessing methods: B-trees, hashing, etc.
• Query languages: SQL, etc.
• User interfaces, forms and reports
• Query processing and query optimization
• Transactions, concurrency control and recovery
• On-line transaction processing (OLTP)
Advanced Database Systems
(mid-1980s–present)
• Advanced data models: extended relational, object-relational, etc.
• Advanced applications: spatial, temporal, multimedia, active, stream and sensor, scientific and engineering, knowledge-based
Advanced Data Analysis: Data Warehousing and Data Mining
(late 1980s–present)
• Data warehouse and OLAP
• Data mining and knowledge discovery: generalization, classification, association, clustering, frequent pattern and structured pattern analysis, outlier analysis, trend and deviation analysis, etc.
• Advanced data mining applications: stream data mining, bio-data mining, time-series analysis, text mining, Web mining, intrusion detection, etc.
• Data mining and society: privacy-preserving data mining
Web-based databas
(1990s–present)
• XML-based databas systems
• Integration with information retrieva
• Data and informatio integration
New Generation of Integrated Data and Information Systems
(present–future)

Road Map

- Motivation
- What is DM?
- DM Tasks
- Applications
- Issues in DM

What motivated DM?

- Necessity is the mother of invention – Plato

- The Explosive Growth of Data:
  - from terabytes to petabytes

  - Data collection and data availability

  - Major sources of abundant data

Contd.....

- Data collection and data availability
- Automated data collection tools, database systems, Web, computerized society
- Major sources of abundant data
- Business: Web, e-commerce, transactions, stocks, ...
- Science: Remote sensing, bioinformatics, scientific simulation, ...
- Society and everyone: news, digital cameras, YouTube

Contd.....

- **We are drowning in data, but starving for knowledge! (Much of the data is never analyzed at all)**

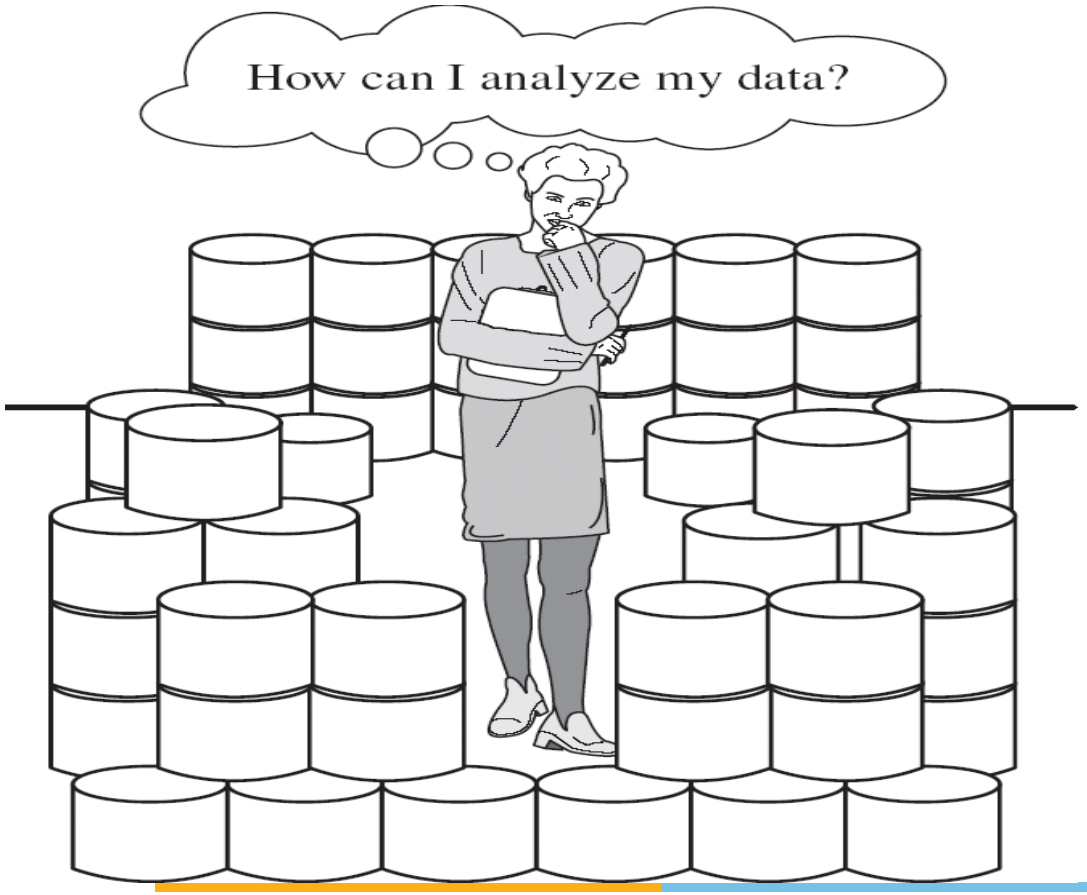

Why DM?

- Consider the data of all educated people ...
- Consider the data of white collar crimes...
- To decide whether the crime rate is increasing or decreasing along with the rate of education.....
- To be more specific , what is the rate of people from rural and urban areas who involve in the crimes.....

Contd .....

- Strategic Decision Making
- Wealth Generation
- Analyzing Trends

Road Map

- Motivation
- What is DM?
- DM Tasks
- Applications
- Issues in DM

What is DM?

- Data mining (knowledge discovery from data)
  - Extraction of interesting (non-trivial, implicit, previously unknown and potentially useful) patterns or knowledge from huge amount of data.

Contd .....

- Alternative names

- Knowledge discovery (mining) in databases (KDD), knowledge extraction, data/pattern analysis, data archeology, data dredging, information harvesting, business intelligence, etc.

DM on what kind of Data?

- Relational Database
- Data Warehouse (is a repository of information collected from multiple sources, stored under a unified schema, and usually resides at a single site)

Contd.....

- Advanced data and information systems
- Object-oriented database
- Temporal DB, Sequence DB and Time serious DB
- Spatial DB
- Text DB and Multimedia DB
- … and WWW

What is DM?

- Data mining (knowledge discovery from data)
  - Extraction of interesting (non-trivial, implicit, previously unknown and potentially useful) patterns or knowledge from huge amount of data.

Types of interestingness

- Frequancy
- Rarity
- Correlation
- Length Of Occurance (for sequance and temporal data)
- Consistency
- Repeating / Periodicity
- Abnormal Behaviour
- Other patterns ……

Contd .....

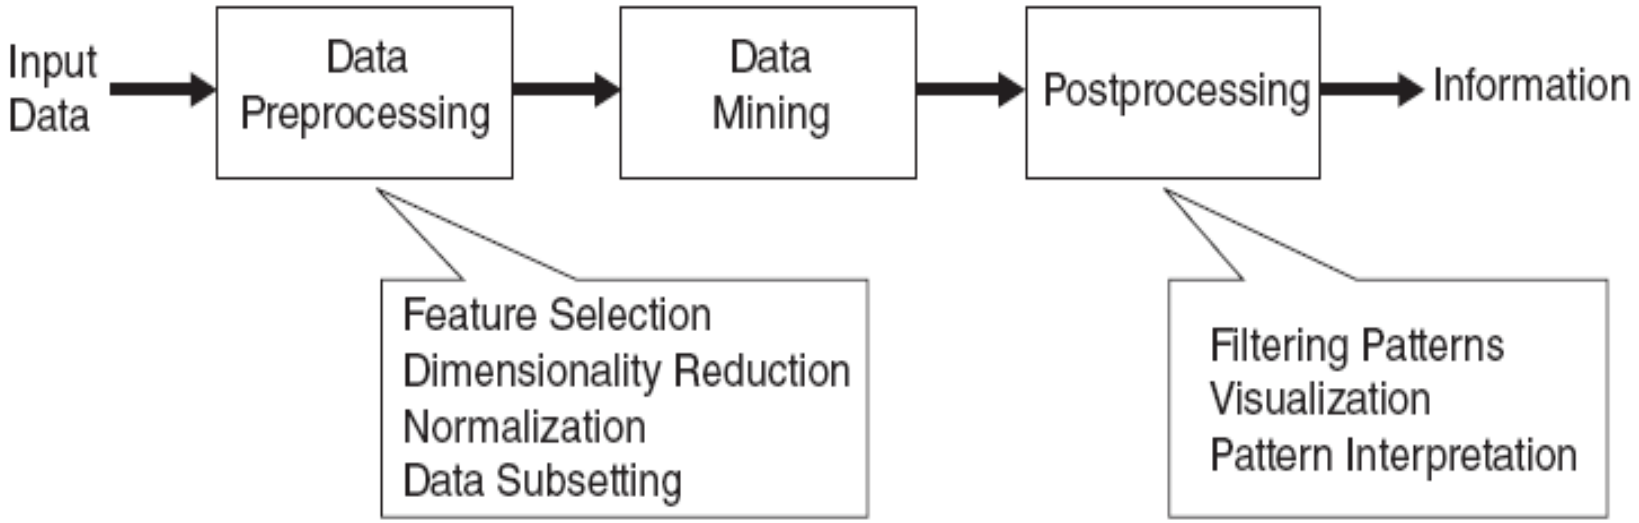

Figure 1.1. The process of knowledge discovery in databases (KDD).

Contd .....

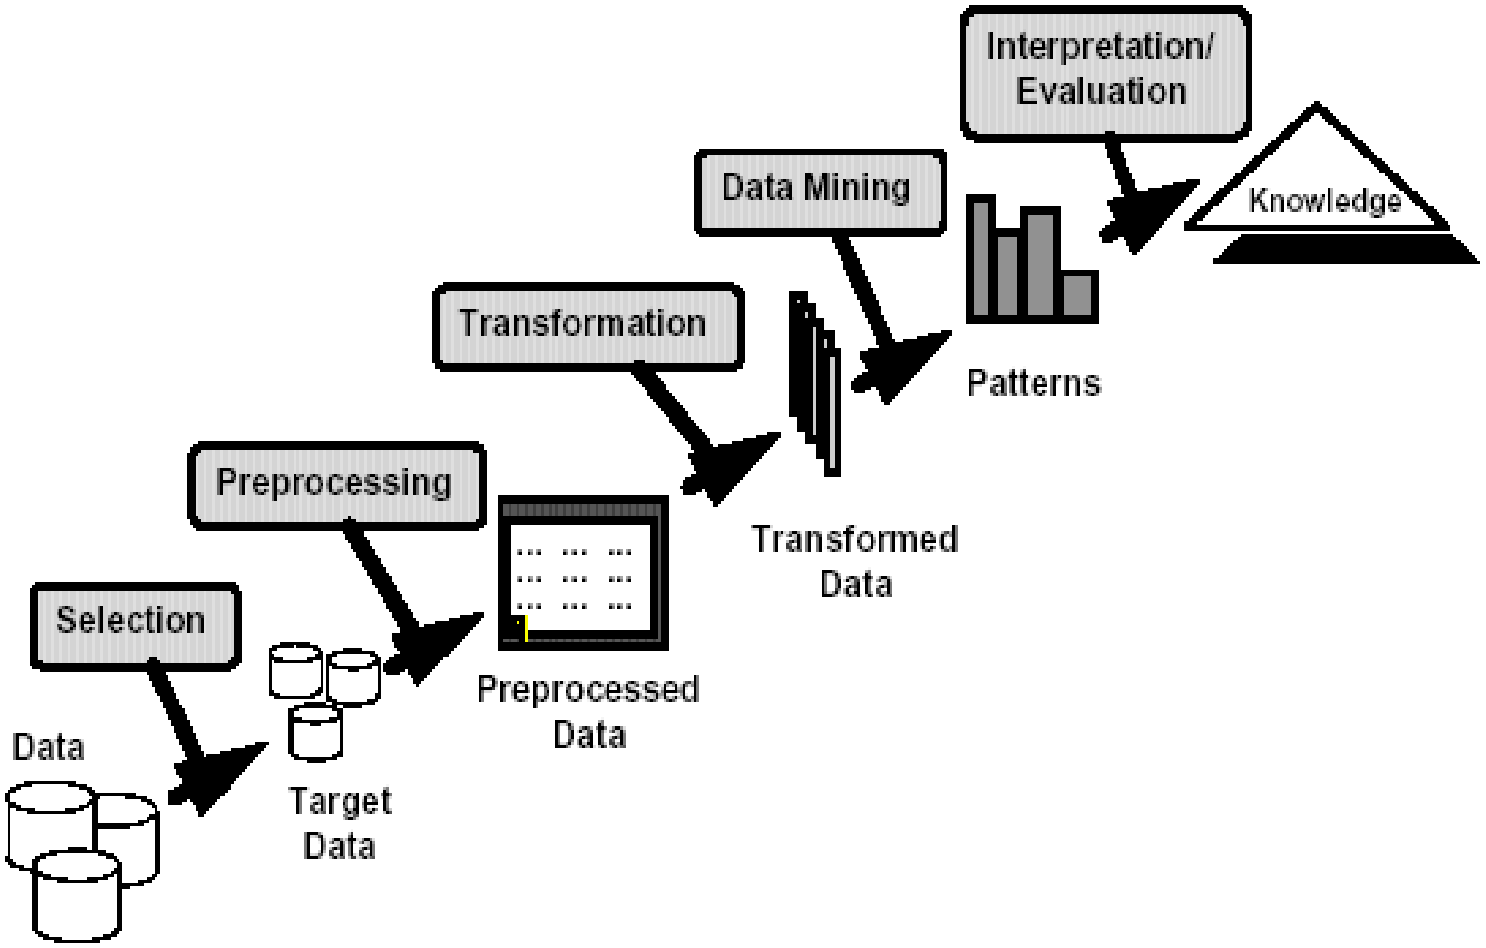
Interpretation/
Evaluation
Knowledge
Data Mining
Patterns
Transformation
Transformed
Data
Preprocessing
Preprocessed
Data
Selection
Target
Data
Data

Contd .....

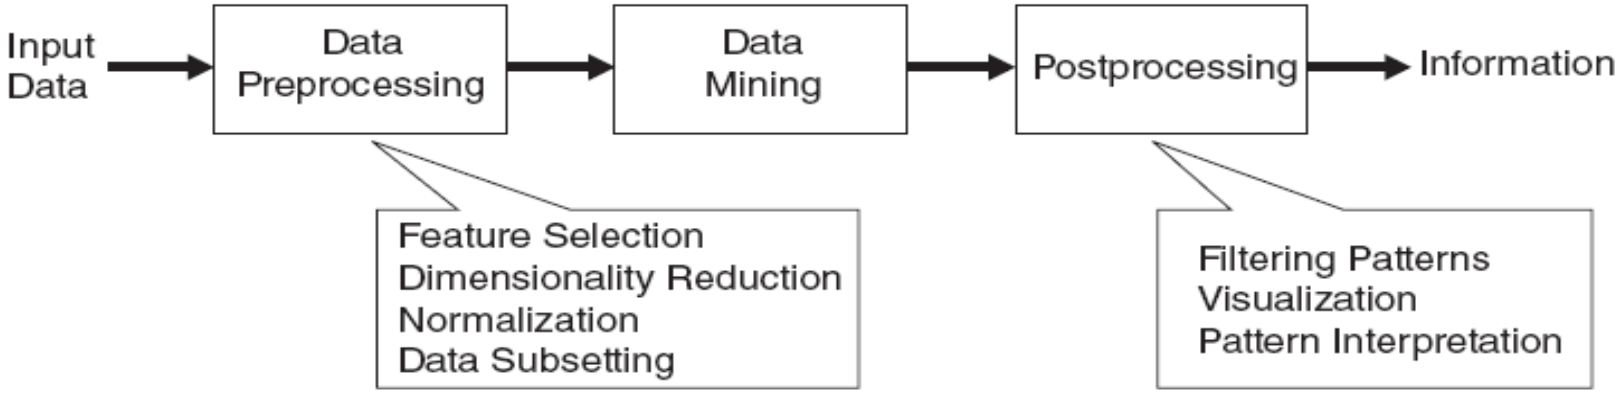

Figure 1.1. The process of knowledge discovery in databases (KDD).

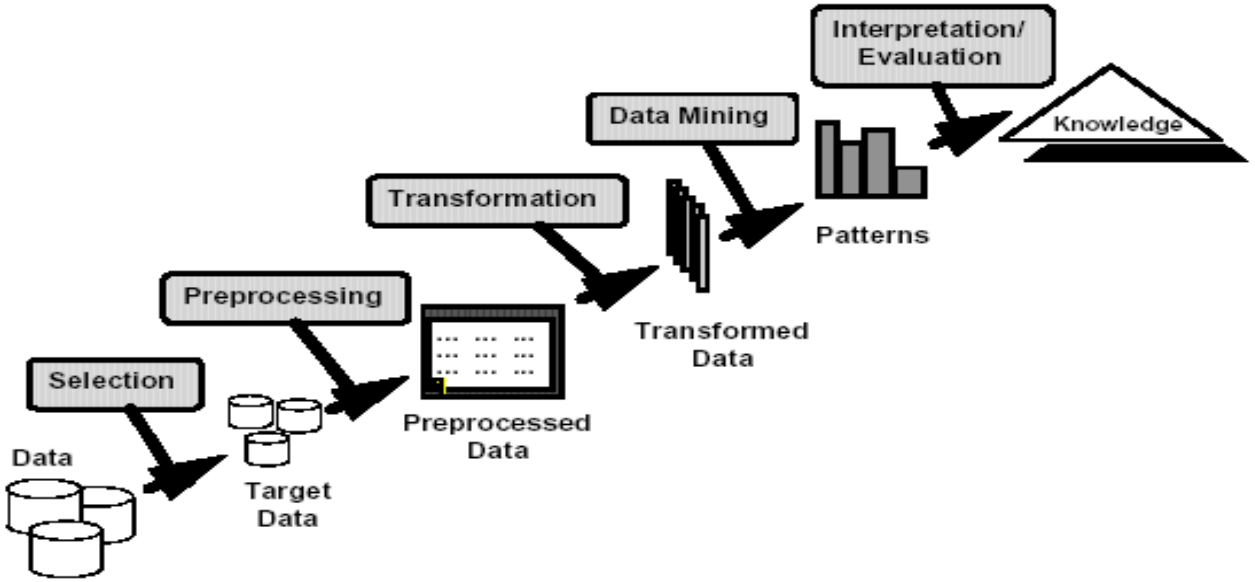

Contd .....

- The process of converting raw data into useful information.
- We always try to represent in a model so that the end user can interpret the reults in a meaningful manner.

What is (not) DM?

- **What is not Data Mining?**
  - Look up phone number in phone directory
  - Query a Web search engine for information about “Amazon”

- **What is Data Mining?**
  - Certain names are more prevalent in certain North India locations (Srivastava Jain, Chawla…)
  - Group together similar documents returned by search engine according to their context (e.g. Amazon rainforest, Amazon.com,)

Data Mining: Confluence of Multiple Disciplines

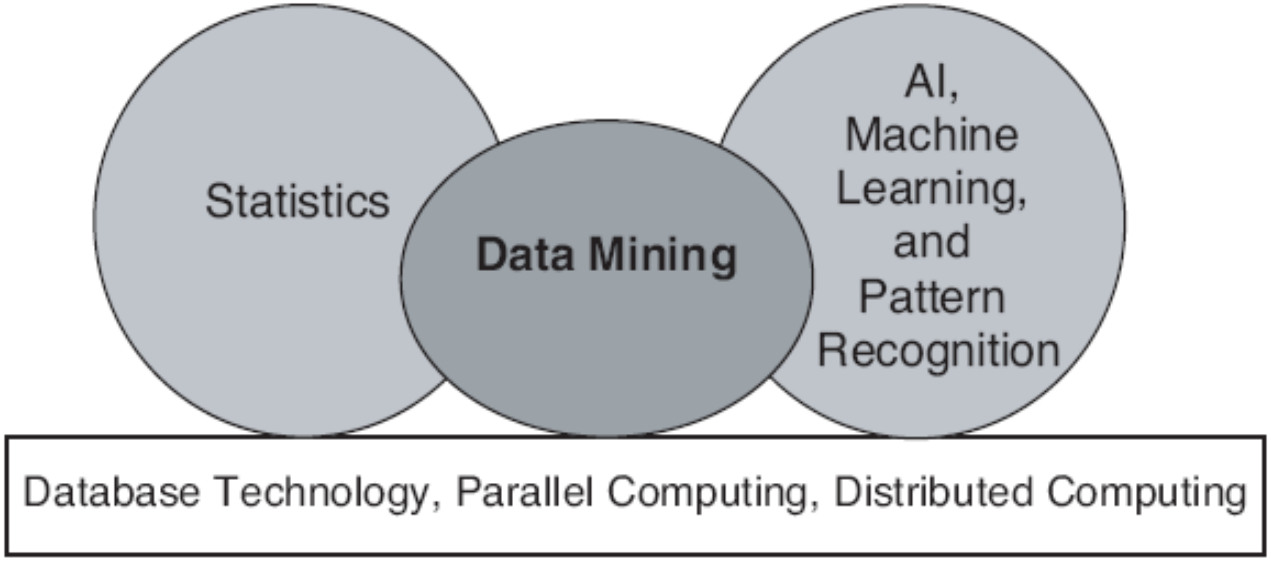

Figure 1.2. Data mining as a confluence of many disciplines.

Road Map

Motivation
What is DM?
DM Tasks
Applications
Issues in DM

DM Tasks

- Prediction Methods
  - Use some variables to predict unknown or future values of other variables.
- Description Methods
  - Find human-interpretable patterns that describe the data.

Contd .....

- Classification [Predictive]
- Clustering [Descriptive]
- Association Rule Discovery [Descriptive]
- Sequential Pattern Discovery [Descriptive]
- Regression [Predictive]
- Deviation Detection [Predictive]

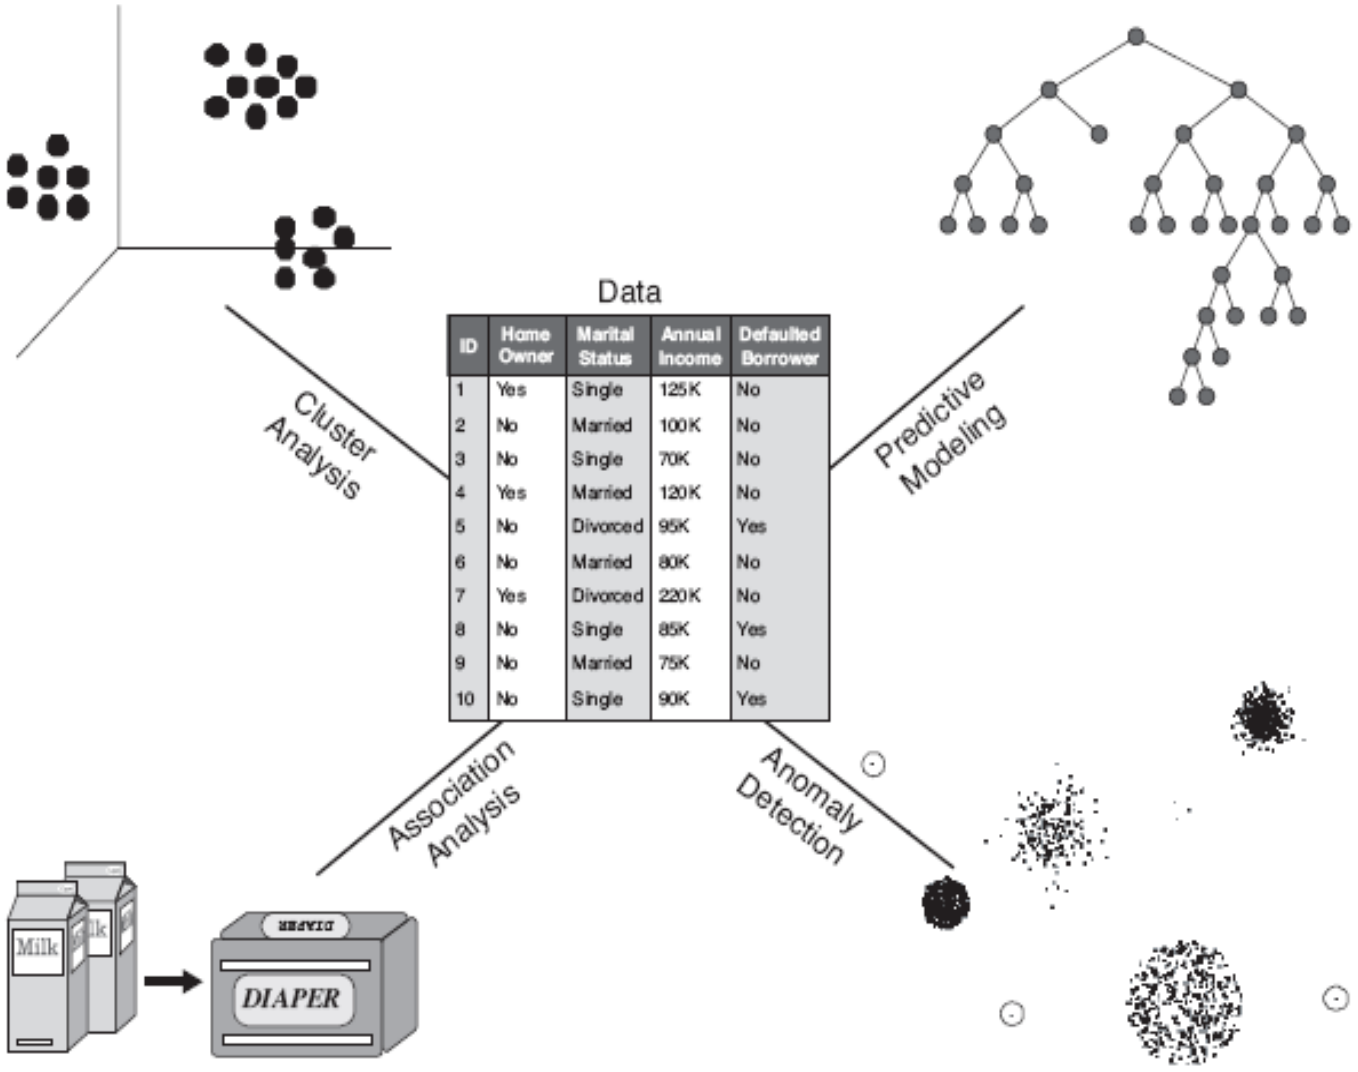

Figure 1.3. Four of the core data mining tasks.

Road Map

- Motivation
- What is DM?
- DM Tasks
- Applications
- Issues in DM

Applications

- Commercial Applications / Business Apllications

  Financial Applications (Bank ,Stock Exchange)
  E-Commerce (Flipcart ,eBay...)

- Scientific Applications

  Astronomical Applications
  Weather Information
  Earth ScienceApplications

- Social Applications

  - FaceBook
  - Economy of People

Classification .....

- **Classification** is the process of finding a MODEL that describes and distinguish data classes or concepts

(a)

age(X, "youth") AND income(X, "high") → class(X, "A")
age(X, "youth") AND income(X, "low") → class(X, "B")
age(X, "middle_aged") → class(X, "C")
age(X, "senior") → class(X, "C")

(b)

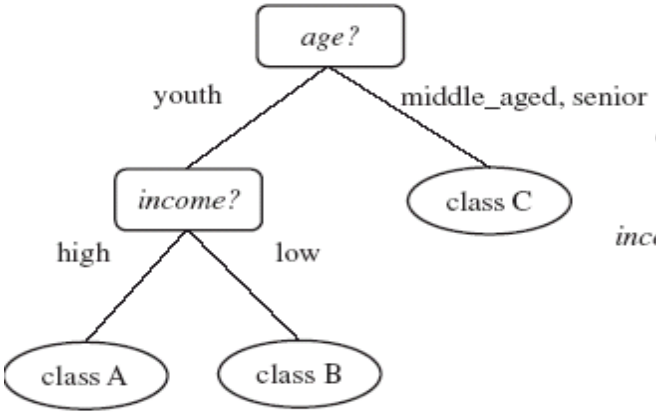

(c)

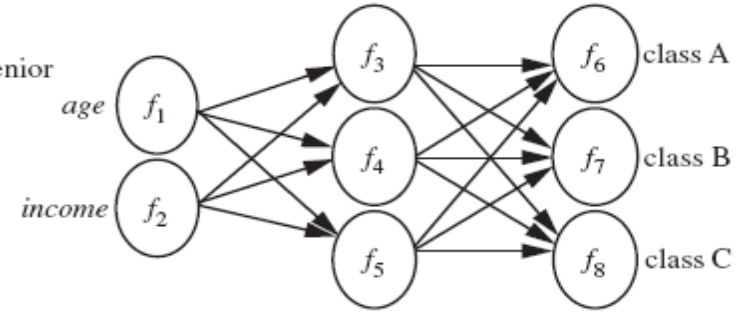

Classification Application 1

- Direct Marketing
  - Goal: Reduce cost of mailing by *targeting* a set of consumers likely to buy a new cell-phone product.
  - Approach:
  - Use the data for a similar product introduced before.
  - We know which customers decided to buy and which decided otherwise. This *{buy, don't buy}* decision forms the *class attribute*.
  - Collect various demographic, lifestyle, and company-interaction related information about all such customers.
  - Type of business, where they stay, how much they earn, etc.
  - Use this information as input attributes to learn a classifier model.

Classification Application 2

- Fraud Detection
  - Goal: Predict fraudulent cases in credit card transactions.
  - Approach:
  - Use credit card transactions and the information on its account-holder as attributes.
  - When does a customer buy, what does he buy, how often he pays on time, etc
  - Label past transactions as fraud or fair transactions. This forms the class attribute.
  - Learn a model for the class of the transactions.
  - Use this model to detect fraud by observing credit card transactions on an account.

Clustering

- In general, the class label are not present in the training data simply they are not known to begin with

- The objects are clustered or grouped based on the principle of *maximizing the intra-cluster similarity* and *minimizing the inter-cluster similarity*

Clustering

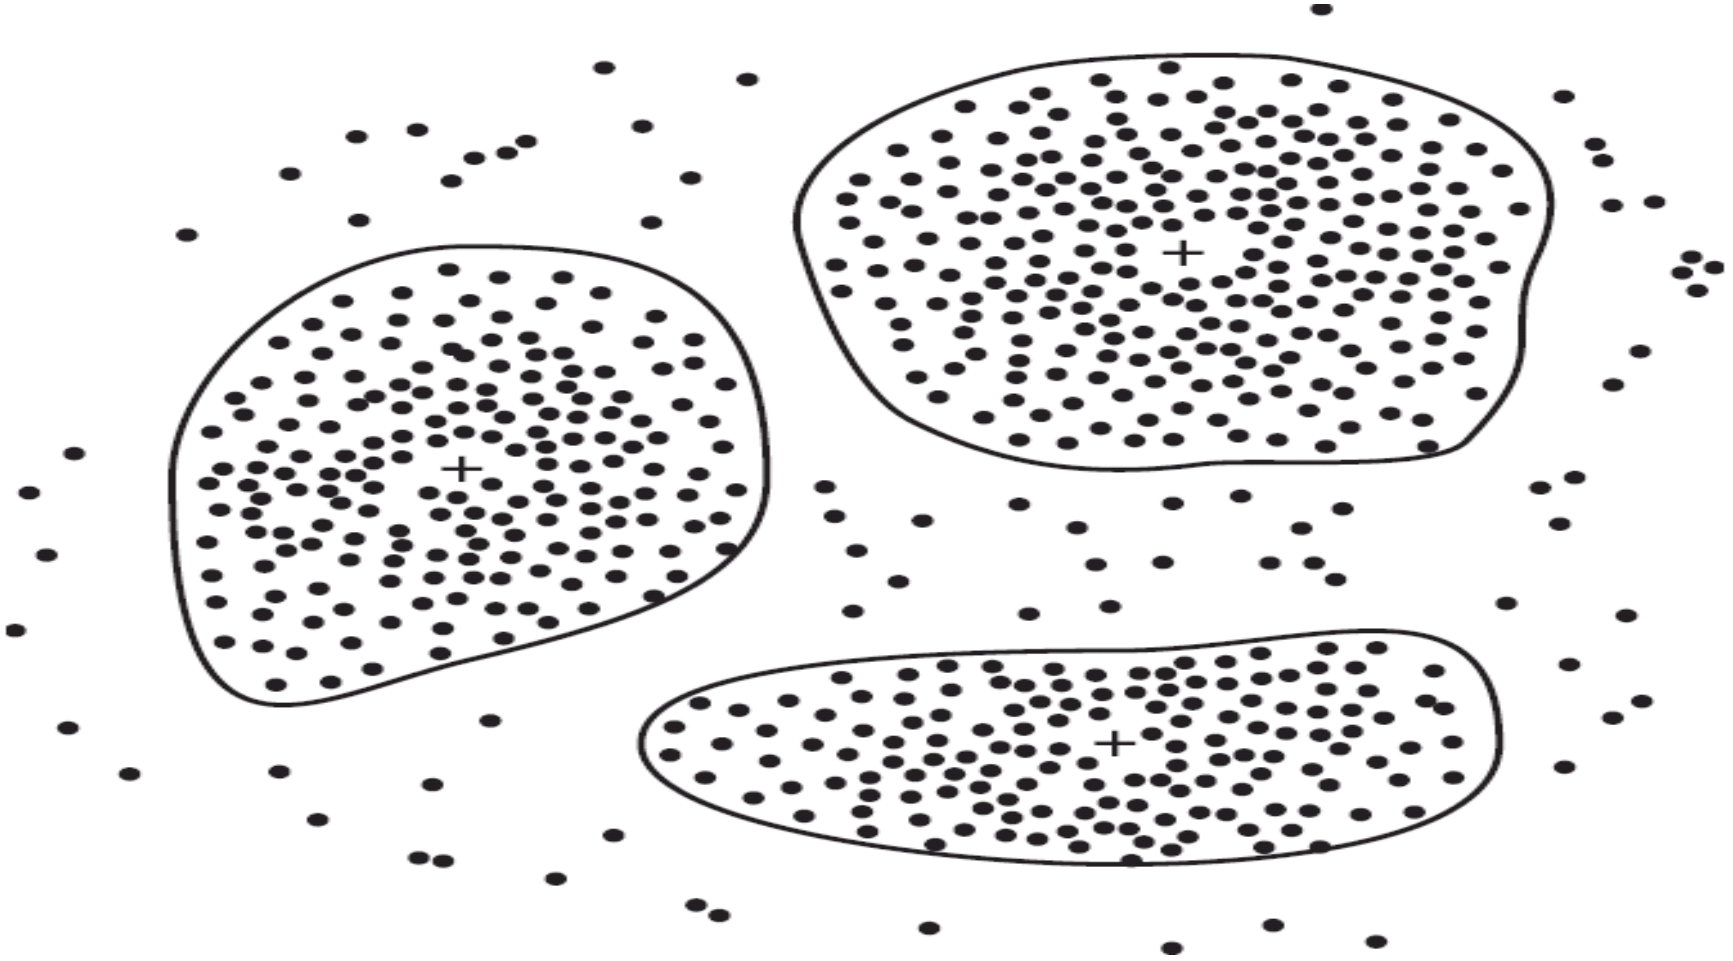

Clustering: Application 1

- Market Segmentation:
  - Goal: subdivide a market into distinct subsets of customers where any subset may conceivably be selected as a market target to be reached with a distinct marketing mix.
  - Approach:
  - Collect different attributes of customers based on their geographical and lifestyle related information.
  - Find clusters of similar customers.
  - Measure the clustering quality by observing buying patterns of customers in same cluster vs. those from different clusters.

Clustering: Application 2

- Document Clustering:
  - Goal: To find groups of documents that are similar to each other based on the important terms appearing in them.
  - Approach: To identify frequently occurring terms in each document. Form a similarity measure based on the frequencies of different terms. Use it to cluster.
  - Gain: Information Retrieval can utilize the clusters to relate a new document or search term to clustered documents.

Association Rule Discovery

- **Frequent patterns** are patterns that occur frequently in data
- Association analysis:
  - Example: buys(X,”computer”) => buys(X,”software”) [support = 1%, confidence = 50%]

Association Rule Discovery: Application 1

- Supermarket shelf management.
  - Goal: To identify items that are bought together by sufficiently many customers.
  - Approach: Process the point-of-sale data collected with barcode scanners to find dependencies among items.
  - A classic rule --
  - If a customer buys diaper and milk, then he is very likely to buy cheese.

Sequential Pattern Discovery

- Applications of sequential pattern mining

  - Customer shopping sequences:
  - First buy computer, then CD-ROM, and then digital camera, within 3 months.
  - Medical treatments, natural disasters (e.g., earthquakes), science & eng. processes, stocks and markets, etc.
  - Telephone calling patterns, Weblog click streams
  - DNA sequences and gene structures

Regression

- Predict a value of a given continuous valued variable based on the values of other variables, assuming a linear or nonlinear model of dependency.
- Greatly studied in statistics, neural network fields.
- Examples:
  - Predicting sales amounts of new product based on advetising expenditure.
  - Predicting wind velocities as a function of temperature, humidity, air pressure, etc.
  - Time series prediction of stock market indices.

Deviation Detection

- Applications:
  - Credit Card Fraud Detection
- Network Intrusion
  - Detection

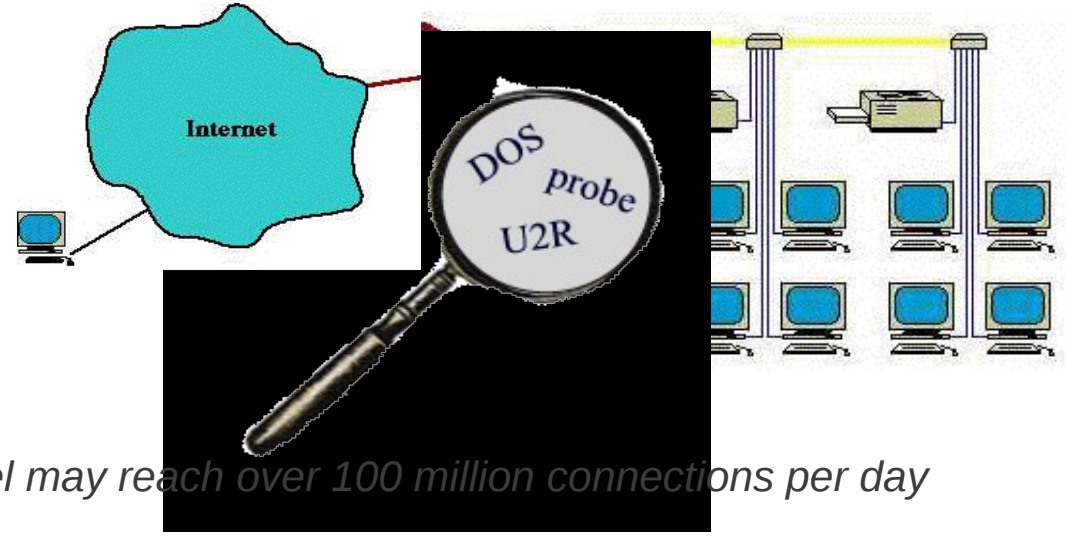

- *Typical network traffic at University level may reach over 100 million connections per day*

Road Map

- Motivation
- What is DM?
- DM Tasks
- Applications
- Issues in DM

- Scalability
- Dimensionality
- Complex and Heterogeneous Data
- Data Quality
- Data Ownership and Distribution
- Privacy Preservation
- Streaming Data

Scalability

- Algorithms generally:
- Operate on data with assumption of in-memory processing of entire data set
- Operate under assumption of developers will address I/O and other performance scaling issues
- Or just don't address scalability within resource constraints at all

Dimensionality

LSST : 1000's of dimensions.

- Massive data stream: ~2Terabytes of image data per
- hour that must be mined in real time (for 10 years).

• Massive 20-Petabyte database: more than 50 billion objects need to be classified, and most will bemonitored for important variations in real time.

• Massive event stream: knowledge extraction in real time for 100,000 events each night.

Complex & Heterogenous data

Data are usually heterogeneous (e.g., databases, images, catalogs, file systems, web interfaces, document libraries, binary, text, structured, unstructured, …)

Data Quality

- Data quality problems are expensive and pervasive
- DQ problems cost hundreds of billion $$$ each year.
- Resolving data quality problems is often the biggest effort in a data mining study.

Data Ownership & Distribution

- **Distributed data are the norm**
- **(across people, institutions,**
- **projects, agencies, nations, …)**

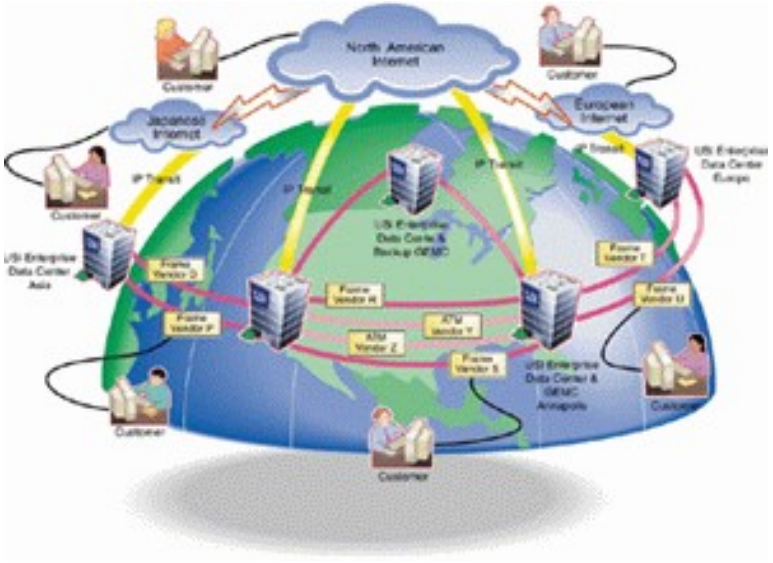

Privacy Preservation

- A Scenario in which two parties owning confidential databases wish to run a data mining algorithm on the union of their databases, without revealing any unnecessary information.

- The need to both protect privileged information and enable its use for research or other purpose.

Streaming Data

- Traditional data mining techniques usually require entire data set to be present.
- Random access (or multiple access) to the data.
- Impractical to store the whole data.
- Simple calculation per data due to time and space constraints.