



BITS Pilani
K K Birla Goa Campus

Operating Systems

Dr. Lucy J. Gudino
Dept. of CS and IS

Last Class



- Segmentation
- Segmentation hardware
- Protection
- Virtual Memory



Virtual Memory

BITS Pilani
K K Birla Goa Campus

Introduction



- How to increase the degree of multiprogramming ?
- Logical vs physical memory
- Virtual memory is a technique that allows the execution of processes that are not completely in memory.
- Motivation:
 - All memory references within a process are logical addresses that are dynamically translated into physical addresses at run time
 - A process may be broken up into a number of pieces (pages or segments) and these pieces need not be contiguously located in main memory during execution.
 - Programs often have code to handle unusual error conditions which is almost never executed.

Contd...



- Arrays, lists, and tables are often allocated more memory than they actually need.
- Certain options and features of a program may be used rarely
- Principle of locality
- trashing is a condition where system spends more time in swapping than executing instructions

Contd...




- Advantages:
 - A program would no longer be constrained by the amount of physical memory that is available
 - Because each user program could take less physical memory, more programs could be run at the same time
 - increases the CPU utilization and throughput
 - Less I/O would be needed to load or swap each user program into memory, so each user program would run faster


Background

- **Virtual memory** – separation of user logical memory from physical memory.
 - Only part of the program needs to be in memory for execution
 - Logical address space can therefore be much larger than physical address space
 - Allows address spaces to be shared by several processes
 - Allows for more efficient process creation

Virtual Memory That is Larger Than Physical Memory



Virtual-address Space




Virtual Memory Implementation



- Virtual memory can be implemented via:
 - Demand paging
 - Demand segmentation

Demand Paging

- Bring a page into memory only when it is needed
 - similar to swapping



Demand Paging

- Bring a page into memory only when it is needed
 - similar to swapping
 - uses **Lazy swapper** – never swaps a page into memory unless page is needed
 - Swapper that deals with pages is a **pager**
- Advantages:
 - Less I/O needed
 - Less memory needed
 - Faster response
 - More users
- Page is needed \Rightarrow reference to it
 - invalid reference \Rightarrow abort
 - not-in-memory \Rightarrow bring to memory

Valid-Invalid Bit


- With each page table entry a valid–invalid bit is associated (**v** \Rightarrow in-memory, **i** \Rightarrow not-in-memory)
- Initially valid–invalid bit is set to **i** on all entries
- Example of a page table snapshot:

Frame #	valid-invalid bit
	v
	v
	v
	v
	i
....	
	i
	i

page table

- During address translation, if valid–invalid bit in page table entry is **i** \Rightarrow page fault


Page Table When Some Pages Are Not in Main Memory



Contd...



If logical address corresponds to page 7?




Page Fault

If there is a reference to a page, first reference to that page will trap to operating system:

page fault

1. Operating system looks at another table to decide:
 - Invalid reference \Rightarrow abort
 - Just not in memory
2. Find empty frame
3. Swap page into frame
4. Reset tables
5. Set validation bit = **v**
6. Restart the instruction that caused the page fault

Steps in Handling a Page Fault



Hardware support

- Same as the hardware for paging and swapping
 - Page table
 - Secondary memory also known as swap device
- How to restart after a page fault?
 - Page fault during instruction fetch
 - Page fault during data fetch

Performance of Demand Paging

- Effective Access Time (EAT)
- $$\text{EAT} = (1 - p) \times \text{ma} + p \times \text{page fault time}$$
- where
- Page Fault Rate p : $0 \leq p \leq 1.0$
 - if $p = 0$ no page faults
 - if $p = 1$, every reference is a fault
- ma : memory access

Sequence of events during page fault



- Trap to the operating system.
- Save the user registers and process state.
- Determine that the interrupt was a page fault.
- Check that the page reference was legal and determine the location of the page on the disk.
- Issue a read from the disk to a free frame:
 - Wait in a queue for this device until the read request is serviced.
 - Wait for the device seek and /or latency time.
 - Begin the transfer of the page to a free frame.

Contd...



- While waiting, allocate the CPU to some other user (CPU scheduling, optional).
- Receive an interrupt from the disk I/O subsystem (I/O completed).
- Save the registers and process state for the other user (if step 6 is executed).
- Determine that the interrupt was from the disk.
- Correct the page table and other tables to show that the desired page is now in memory.
- Wait for the CPU to be allocated to this process again.
- Restore the user registers, process state, and new page table, and then resume the interrupted instruction.

Demand Paging Example

Memory access time = 200 nanoseconds

Average page-fault service time = 8 milliseconds

One access out of 1,000 causes a page fault

$EAT = (1 - p) \times ma + p \times \text{page fault time}$

$EAT = (1 - p) \times 200\text{ns} + p \times (8 \text{ ms})$

then

$EAT = 8.2 \text{ microseconds.}$

This is a slowdown by a factor of 40!!