DATA STORAGE TECHNOLOGIES & NETWORKS

(CS C446, CS F446 & IS C446)

LECTURE 26– STORAGE

vnode

- Information stored in vnode
  - flags (used for locking vnode & identifying generic attributes)
  - Various reference counts
    - Number of file entries that are open for read and/or write that references vnode
    - Number of file entries that are open for writing that reference the vnode
    - Number of pages and buffers that are associated with the vnode
  - A pointer to the mount structure describes the file system that contains the object represents by the vnode
  - Various information is used to do file read ahead
  - A reference to an NFS lease
  - A reference to state about special devices, sockets, and FIFOs
  - Pointer to the set of vnode operations defined for the object
  - Pointer to private information needed for the underlying object (for local file system, this pointer will reference an inode and for NFS it is an nfsnode)
  - The type of the underlying object (regular file, character / block device, directory etc..)
  - Clean and dirty buffers associated with vnode
  - A count (number of buffer write operation in progress)

In-Memory File System Structures

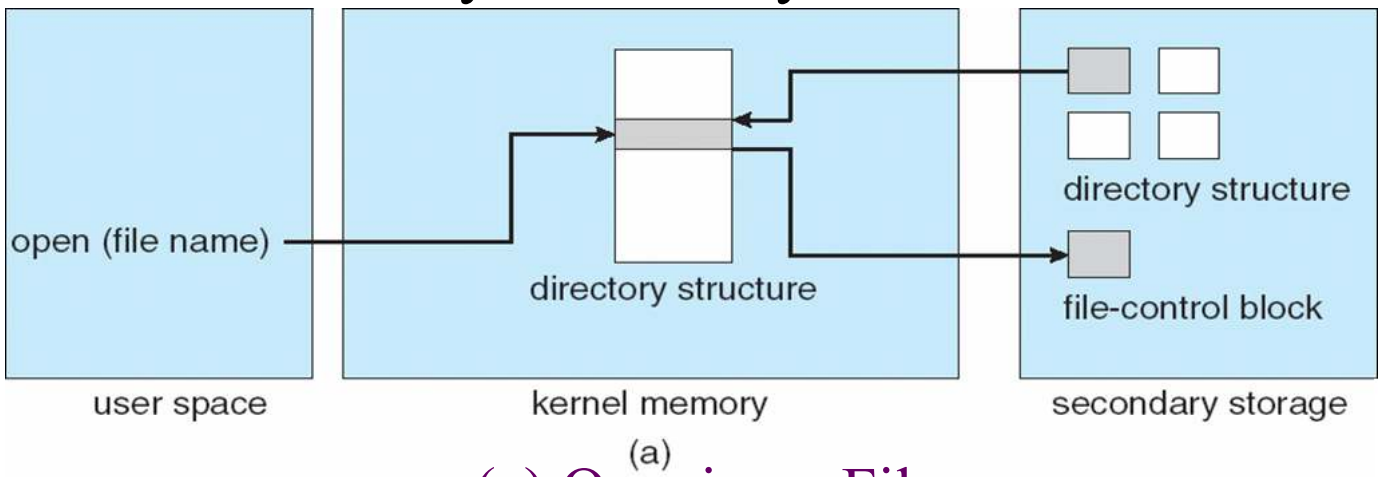

(a) Opening a File

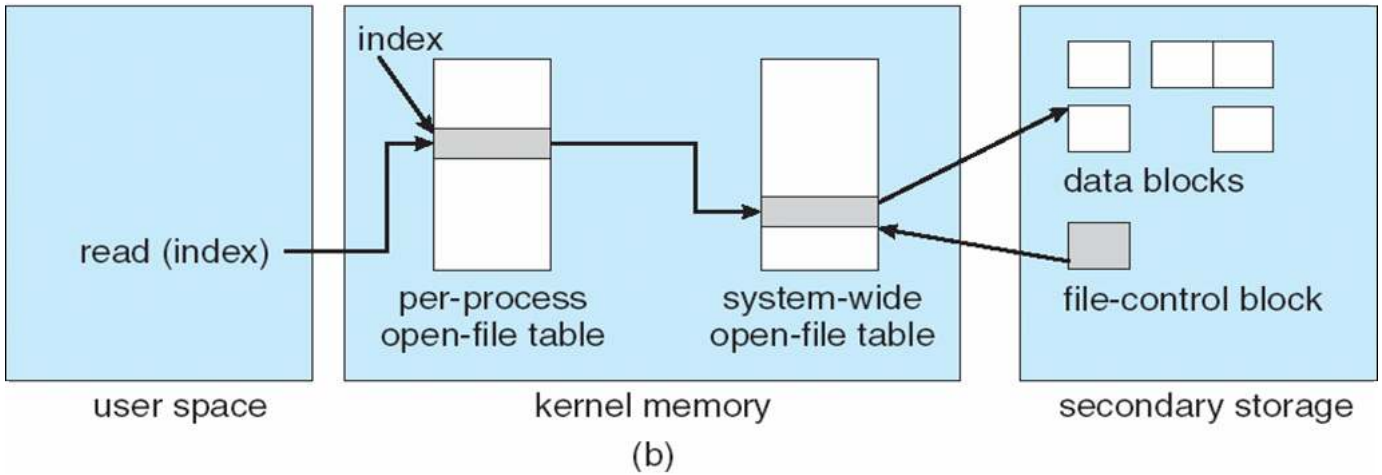

(b)Reading a File

File descriptors, file table and inode table

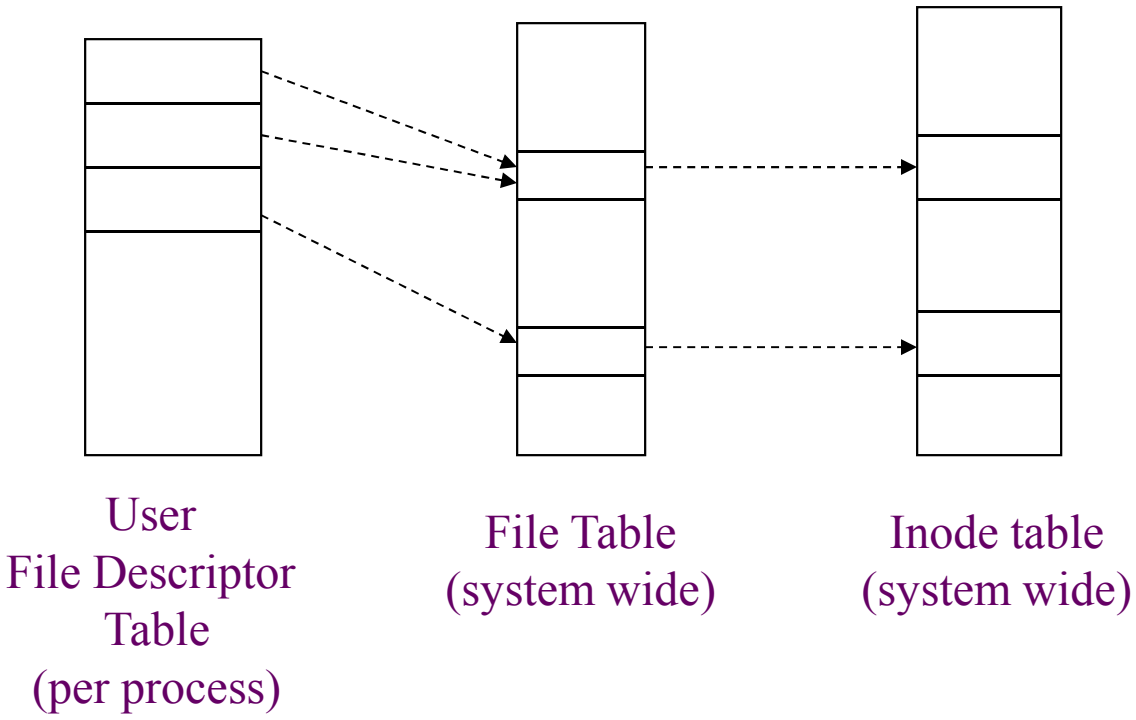

File system layout

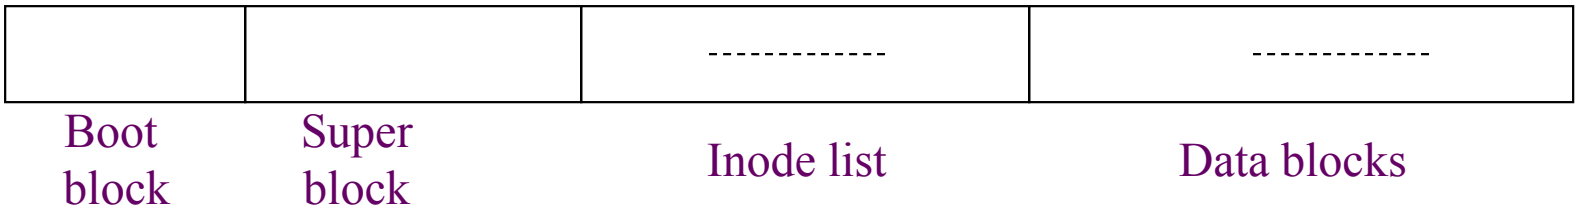

- Boot block occupies the beginning of a file system
  - Typically first block and may contain boot strap code.
  - Only one boot block is needed to boot the system, but every file system has a (possibly empty) boot block.
- Super block describes the state of a file system.
  - How large it is
  - How many files it can store
  - Where to find free space on the file system and
  - Other information.

- Inode list is the list of inodes that follow the super block in the file system
  - Administrators specify the size of the inode list when configuring a file system.
  - Kernel reference inode by index into the inode list.
  - One inode is root inode (by which the directory structure of file system is accessible after execution of mount system call)
- Data blocks starts at the end of inode list and contain file data and administrative data
  - An allocated data block can belong to one and only one file in the file system.

File System Implementation

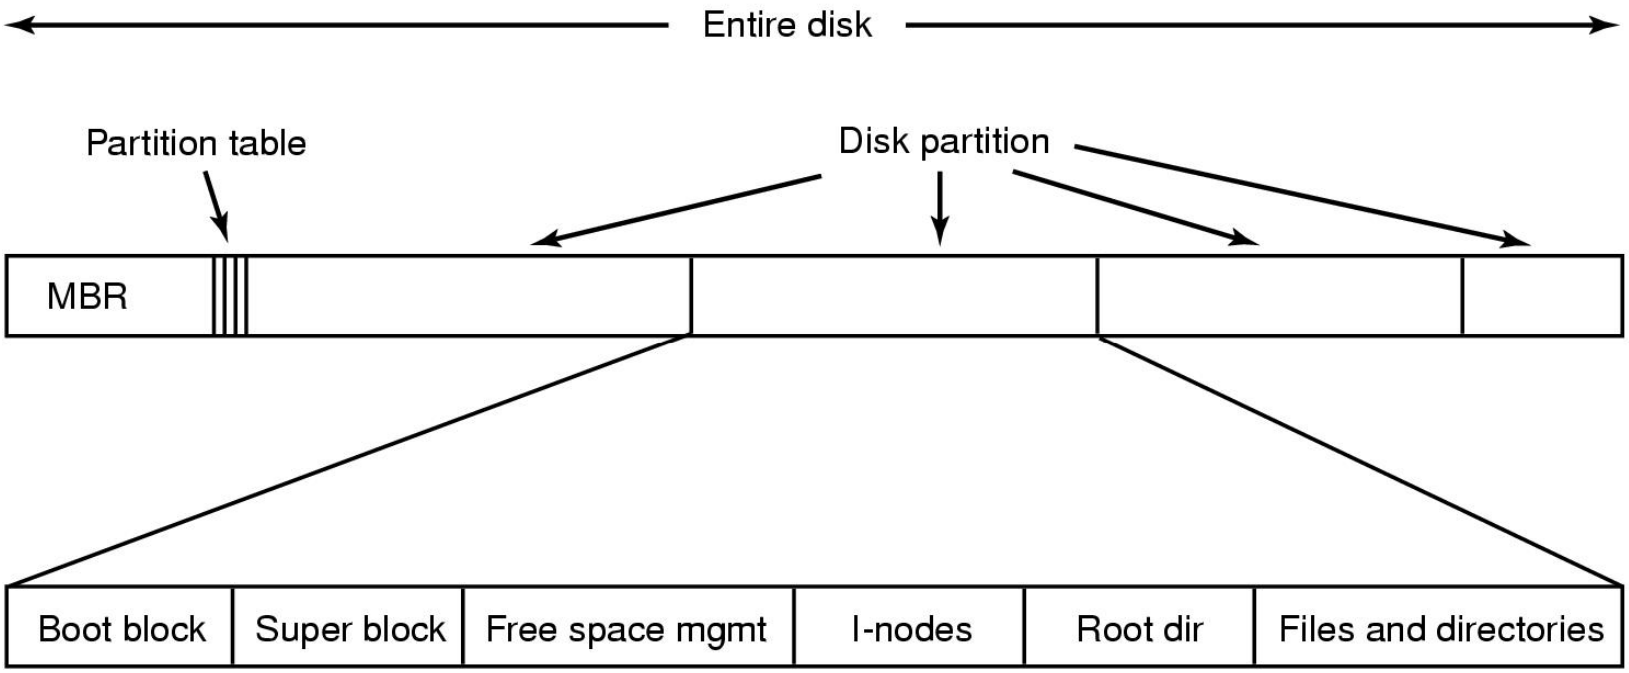

A possible file system layout

Virtual File Systems

- Virtual File Systems (VFS) provide an object-oriented way of implementing file systems.
- VFS allows the same system call interface (the API) to be used for different types of file systems.
- The API is to the VFS interface, rather than any specific type of file system.

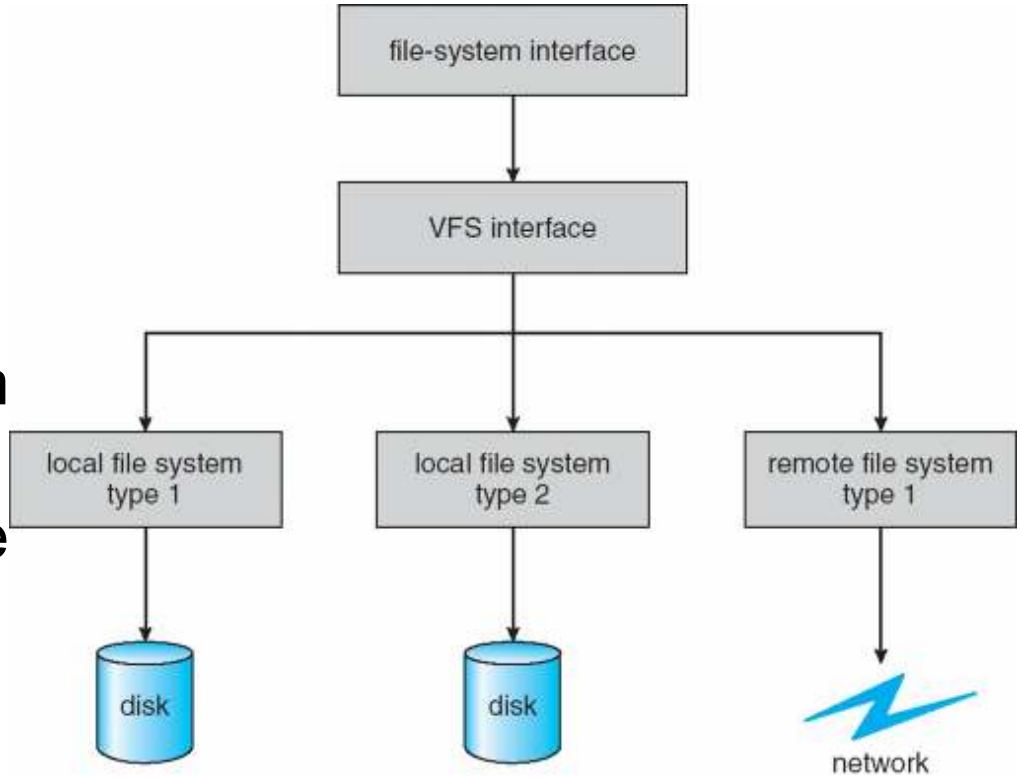

Buffering

- Buffer Cache
  - Memory buffer for data being transferred to and from disks
  - Cache for recently used blocks
    - 85% hit rate is typical
    - Typical buffer size = 64KB virtual memory
    - Buffer pool – hundreds of buffers
  - Consistency issue
    - Each disk block mapped to at most one buffer
    - Buffers have dirty bits associated
    - When a new buffer is allocated and if its disk blocks overlap with that of an existing buffer, then the old buffer must be purged.

Buffer Pool Management

- Buffer pool is maintained as a (separately chained) hash table indexed by a buffer id
- The buffers in the pool are also in one of four lists:
  - Locked list:
    - buffers that are currently used for I/O and therefore locked and cannot be released until operation is complete
  - LRU list:
    - A queue of buffers – a recently used item is added at the rear of the queue and when a buffer is needed one at the front of the queue is replaced.
    - Buffers staying in this queue long enough are migrated to an Aged list
  - Aged List:
    - Maintained as a list and any element may be used for replacement.
  - Empty List
- When a new buffer is needed check in the following order:
  - Empty List, Aged List, LRU list

The buffer cache

- Data access from a file
  - Process request to the kernel
  - Kernel brings data and auxiliary information (super block, inode etc)it into the memory
  - Process examines it
  - Process can alter / read / store it into sec memory back
- Kernel uses internal data buffers
  - Reduces frequency of disk access
  - System response time and throughput will be better
  - Contains the data in recently used disk blocks.
  - It is called buffer cache (a software structure, not hardware).
  - Places between file subsystem block and device driver block
  - Improves the write performance by accumulating multiple writes in the same block before writing it into disk.